

Building a Brand AI Can Find: Mentions, Consistency, and Reputation – Fatma Yaren Birben

A Working Framework for Earning Citations Across ChatGPT, Perplexity, Claude, and Gemini

In this guide, you'll get:

- The three-pillar framework AI systems actually use to decide which brands to mention and recommend
- A self-scoring Brand AI Footprint Audit to find your specific gaps
- A measurement template for tracking AI-referred visibility, traffic, and revenue

"Be the answer. On every search. Every AI. Every time."

1. Executive Summary & Problem Statement

For two decades, online visibility meant one thing: rank on page one of Google. That equation no longer holds by itself. A growing share of buyer research now happens inside an AI system that synthesizes an answer directly — ChatGPT, Perplexity, Claude, Gemini — instead of returning a ranked list of links for the buyer to click through.

This isn't a future problem. FlyRank's own audit data already shows the pattern: brands ranking on page one of Google for their core terms frequently show up nowhere when the same question is put to Perplexity or Claude. Strong search rankings and strong AI visibility are related, but they aren't the same system, and optimizing for one doesn't guarantee results in the other.

The mechanics explain why. Traditional search is a retrieval problem — match a query to the most relevant indexed page and rank it. Generative AI answers are a synthesis problem — the model draws on patterns learned and retrieved across many sources, most of them unstructured, and assembles a single answer. That answer typically names one brand, occasionally three, rarely ten. There is no page two.

The cost of this gap isn't only visibility — it's measurement. Audits across roughly 200 brands run through FlyRank's free AI Visibility tool found that AI-referred revenue and conversions are routinely under-attributed in standard GA4 setups, and where AI traffic is tracked at all, it converts at meaningfully different rates than organic search traffic. Brands are frequently losing ground on a channel they don't even know they're being measured on.

None of this makes traditional SEO obsolete. It means a second, structurally different discovery system now sits alongside it, with its own rules for what gets found, trusted, and recommended.

The rest of this guide breaks that system into three pillars — mentions, consistency, and reputation — and gives you a way to score your brand against each one.

Three questions every brand should be able to answer:

1. Can you find your brand mentioned somewhere you don't own or control?
2. Would an AI system describe your product correctly and specifically, without you prompting it first?
3. If it did, is there a third-party source that backs that description up?

If you hesitated on any of the three, you already know which pillar to start with.

2. The 3 Pillars of AI Brand Visibility

These three pillars aren't abstract theory — they're the same framework FlyRank uses when auditing a brand's AI visibility, because they map to the questions a model effectively has to resolve before it will name a brand in an answer: Have I seen this brand mentioned, and where? Do the facts about it agree everywhere I've encountered them? And is anyone besides the brand itself vouching for it?

2.1 Pillar 1 — Mentions & Citations Across the Web

Traditional SEO trained an entire industry to think in links. A backlink passed authority; an unlinked brand name in a paragraph was, at best, a nice-to-have. That hierarchy breaks down for LLMs.

Language models are trained on and retrieve from enormous volumes of text, and most of that text has no hyperlinks in it at all — forum threads, comparison articles, review pages, transcripts, PDFs. When your brand name shows up in that text alongside your category, a competitor, or a specific use case, that's a usable signal whether or not it's clickable. FlyRank's visibility work is built around exactly this shift: the target isn't just search-engine indexing, it's being cited — showing up as a source an AI system pulls from and repeats.

This changes what "authority" means in practice. Volume and diversity of mention now sit alongside link equity as a distinct signal. A brand named consistently across a wide spread of independent sources — reviews, comparisons, forums, press — builds a stronger association than a brand with a handful of high-authority backlinks and thin mention volume everywhere else.

These mentions accumulate across a fairly predictable set of surfaces:

- Owned content that names your product and category explicitly and in plain language — not just brand voice, but literal, extractable description

- Earned press and guest content that references your brand in an independent context
- Review and ratings platforms — G2, Capterra, Trustpilot, and category-specific sites
- Comparison and "best of" content — the roundup and listicle format LLMs draw on heavily for recommendation-style queries
- Community discussion on Reddit, forums, and Q&A platforms where your category actually gets debated
- Bing and Bing Webmaster Tools, since Bing indexing feeds directly into Microsoft's AI surfaces — worth tracking combined search-and-AI discovery data here specifically, not just organic clicks

At FlyRank, citation monitoring is a standing part of every client audit, not a one-time push, because mention presence shifts constantly as the web — and the models trained on it — keep updating.

Callout: Mention Sources to Track

Source Type	Why It Matters	How to Check
Owned content (blog, comparison pages)	Sets your own language for your category, in your control	Audit whether key pages name your product and category explicitly, in plain terms
Earned press & guest articles	Independent framing an AI system can cross-reference against your own claims	Search brand name + category term; review coverage from the last 12 months
Review platforms (G2, Capterra, Trustpilot, etc.)	High-trust, frequently crawled, already structured for comparison	Confirm profiles are claimed, current, and complete
Comparison & "best of" roundups	Feeds recommendation-style AI queries directly	Search "[category] + best / top / alternatives"; check whether you appear
Community discussion (Reddit, forums, Q&A)	Heavily weighted for research-style and comparison queries	Search brand + category on Reddit/forums; note accuracy and sentiment

Source Type	Why It Matters	How to Check
Bing / Bing Webmaster Tools	Feeds Microsoft's AI surfaces, including Copilot, directly	Review the AI Performance report inside Bing Webmaster Tools

2.2 Pillar 2 — Entity Facts & Multi-Platform Consistency

An AI system doesn't read your website the way a person does. It builds something closer to an entity profile: a set of facts it associates with your brand name — category, offering, audience, differentiation — assembled from every instance of your brand it has encountered. Agreement across those instances functions as a confidence signal.

This is what makes consistency a technical requirement rather than a branding nicety. If your homepage describes you one way, your G2 listing describes you slightly differently, and a press mention from two years ago describes you a third way, the model has to resolve that disagreement somehow. It typically resolves it by hedging, or by defaulting to whichever source it judges more reliable — which may not be you. The brands that show up cleanly and specifically in AI answers are consistently the ones whose facts agree with themselves everywhere, not just on their own site.

Three mechanics do most of the work here:

- **Structured data (schema markup).** Organization, Product, FAQ, Article, and Review schema give machines an explicit, unambiguous version of your facts instead of forcing an AI system to infer them from prose. FlyRank's visibility framework treats this as foundational: content has to be structured for machine extraction, not just written for a human reader.
- **llms.txt.** A newer, plain-text convention that gives AI crawlers a direct, curated summary of what your site contains and where the pages that matter most actually live — a purpose-built alternative to making a model guess its way through your full sitemap.
- **robots.txt and crawler access.** None of the above matters if AI crawlers can't reach your pages. A blocked or misconfigured robots.txt can quietly exclude a brand from AI retrieval entirely, regardless of how strong the content itself is.

FlyRank's own AI Visibility audit checks exactly this stack on every site it evaluates — homepage, robots.txt, llms.txt, sitemap, and a sample of key pages — because inconsistency or blocked access at any single point undermines everything built around it.

Callout: Entity Consistency Checkpoints

Checkpoint	Where to Check	Common Failure Mode
Organization schema	Page source / JSON-LD, schema testing tools	Missing entirely, or present but outdated (old name, old category)
Product/Service schema	Key product or service pages	Naming is inconsistent between product pages and marketing pages
FAQ schema	High-intent pages — pricing, comparison, features	FAQ content is visible on the page but not marked up for machines
llms.txt	yourdomain.com/llms.txt	File doesn't exist, or exists but is stale or incomplete
robots.txt	yourdomain.com/robots.txt	AI crawlers (GPTBot, ClaudeBot, PerplexityBot, etc.) are unintentionally blocked
Cross-source fact agreement	Homepage vs. review profiles vs. press vs. directories	Category, positioning, or core facts differ meaningfully by source

2.3 Pillar 3 — Third-Party Reputation & Authority Signals

A brand's own website is, by definition, the least objective source of information about that brand. AI systems are effectively trained to account for this. Independent corroboration — a third party describing your product, comparing it to competitors, or vouching for outcomes — carries more evidentiary weight than the identical claim made on your own homepage, because it doesn't carry the same incentive problem.

This is the same underlying logic that made backlinks valuable in traditional SEO, but the mechanism has shifted. A backlink told a search engine a page was worth crawling. A third-party mention or citation tells an AI system that a claim about the brand has independent support — which matters even absent a link entirely. The two signals overlap (authoritative domains still matter to both) but they diverge in what they're actually doing: link equity governs crawl priority

and ranking weight; AI-citation value governs whether a claim is trusted enough to be repeated in a generated answer.

The signal categories that build this kind of reputation over time:

- Independent press coverage — journalism and industry publications discussing your company, product, or category position
- Structured reviews — aggregated third-party sentiment on platforms like G2, Capterra, and Trustpilot
- Comparison content you don't control — third-party "X vs. Y" pieces and review roundups
- Backlinks and citations from topically authoritative domains — sites already trusted specifically within your category, not just high-traffic sites generally
- Partner and integration mentions — case studies, co-marketing pages, and integration directories that reference you in a real-world context

FlyRank's client results reflect this directly: the pattern isn't just improved rankings — it's brands becoming the cited answer across both Google and AI systems like ChatGPT simultaneously, which only happens once independent sources are reinforcing the same facts a brand is telling about itself.

Callout: Reputation Signal Types

Signal Type	Example Channels*	What It Signals to an AI System
Independent press	Trade publications, industry news sites	Third-party validation of category position or specific claims
Structured reviews	G2, Capterra, Trustpilot	Aggregated, quantifiable sentiment from users, not the brand
Comparison content	Third-party "X vs. Y" articles, review roundups	Competitive framing you don't control, which reads as more neutral
Authoritative backlinks/citations	Category-relevant sites with established trust	Reinforcement from sources already treated as reliable in

Signal Type	Example Channels*	What It Signals to an AI System
		your space
Partner & integration mentions	Case studies, co-marketing pages, integration directories	Real-world corroboration of use cases and outcomes

**Example channels are illustrative categories, not a prescriptive or exhaustive list — the right mix depends on your category and where your buyers already look for validation.*

3. Practical Implementation: Brand AI Footprint Audit Checklist & Remediation Scorecard

The three pillars give you a way to think about AI visibility. This section turns that thinking into a number. Score your brand against all three pillars using the checklist below, and you'll know exactly where you stand and what to fix first — not as a one-time snapshot, but as a benchmark you can re-run every quarter as your footprint changes.

How to Score

Score each item below 0, 1, or 2:

- **0 — Absent.** Not implemented, or you can't verify it exists when you check.
- **1 — Partial.** In place inconsistently, outdated, or incomplete.
- **2 — Fully Implemented.** Current, complete, and verified everywhere you checked.

Add your scores across all 20 items (40 points possible) and divide by 40 to get a percentage. That places your brand in one of three tiers — the same logic behind FlyRank's own Site Readiness Score, used across client audits to turn a diagnostic into a single number leadership can track over time.

Tier	Score Range	What It Means
Foundational (0–40%)	0–16 points	The basics aren't in place yet. AI systems have little reliable material to work with.

Tier	Score Range	What It Means
Developing (41–75%)	17–30 points	Meaningful groundwork exists, but gaps are creating inconsistency or blind spots.
Advanced (76–100%)	31–40 points	A strong footprint across all three pillars. Focus shifts to maintenance and expansion.

A brand scoring 24 out of 40 (60%), for example, sits in the Developing tier — real progress on mentions and structure, likely with gaps in the reputation layer still capping visibility.

Part A — Mentions & Citations (Items 1–7)

#	Checklist Item	How to Verify	Remediation Action	Priority	Score (0–2)
1	Brand surfaces in AI answers for core category queries	Run 10–15 branded/category prompts across ChatGPT, Perplexity, Claude, Gemini; log results	Build comparison content; pursue press targeting category queries	High	
2	Owned content explicitly names product/category in plain language	Audit top 10–20 pages for explicit, extractable descriptions	Rewrite vague or jargon-heavy copy into plain factual statements	High	
3	Active	Search	Pitch	High	

#	Checklist Item	How to Verify	Remediation Action	Priority	Score (0–2)
	presence in comparison and "best of" content	"[category] + best/top/alternatives"; check inclusion	inclusion; publish original comparison content		
4	Review platform profiles are claimed and complete	Check G2, Capterra, Trustpilot, and category-specific sites	Claim, complete, and actively request reviews	High	
5	Earned press coverage exists and is current	Search brand + category term; review coverage from the last 12 months	Pursue PR and guest content targeting category-defining queries	Medium	
6	Brand appears accurately in community discussion	Search brand + category on Reddit and relevant forums	Engage authentically; correct inaccurate claims where appropriate	Medium	
7	Brand is indexed and performing in Bing	Check the AI Performance report in Bing Webmaster Tools	Submit sitemap; resolve crawl errors	Medium	

Part B — Entity Consistency (Items 8–14)

#	Checklist Item	How to Verify	Remediation Action	Priority	Score (0–2)
8	Organization schema is implemented and current	Check JSON-LD via page source or a schema testing tool	Implement or update Organization schema with accurate fields	High	
9	robots.txt allows AI crawler access	Check yourdomain.com/robots.txt for GPTBot, ClaudeBot, PerplexityBot, etc.	Remove unintended blocks; confirm crawl access	High	
10	llms.txt is present and accurate	Check yourdomain.com/llms.txt	Draft and publish a current llms.txt	High	
11	Core facts agree across owned and third-party sources	Compare homepage vs. G2/press/directories for category, positioning, claims	Standardize a single source-of-truth fact sheet; update all listings	High	
12	Product/Service schema is implemented and consistent	Check schema across all product or service pages	Standardize naming; implement schema site-wide	Medium	
13	FAQ schema	Check	Add FAQ	Medium	

#	Checklist Item	How to Verify	Remediation Action	Priority	Score (0–2)
	is implemented on high-intent pages	pricing, comparison, and feature pages for FAQ markup	schema to existing FAQ content		
14	Sitemap is current and complete	Check yourdomain.com/sitemap.xml against the live site	Regenerate and resubmit sitemap	Medium	

Part C — Reputation & Authority (Items 15–20)

#	Checklist Item	How to Verify	Remediation Action	Priority	Score (0–2)
15	Brand appears in third-party comparison content	Search "[brand] vs. [competitor]"	Identify gaps; pitch inclusion or publish a comparison response	High	
16	Backlinks come from topically authoritative domains	Backlink audit filtered by topical relevance, not just domain authority	Prioritize outreach to relevant, trusted sites over high-DA/low-fit sites	Medium	
17	Case studies are published and	Confirm case studies exist, are indexed,	Publish or restructure case studies	Medium	

#	Checklist Item	How to Verify	Remediation Action	Priority	Score (0–2)
	structured for extraction	and state outcomes explicitly	with clear, quotable outcomes		
18	Active, optimized presence on category review platforms	Search "[brand] reviews"	Claim and optimize profiles; request reviews on a regular cadence	Medium	
19	Independent press coverage exists within the last 12 months	Search brand + category in news and trade press	Pursue PR outreach targeting category-relevant publications	Medium	
20	Partner and integration pages reference the brand accurately	Review partner directories, integration marketplaces, co-marketing pages	Update outdated references; pursue new partner mentions	Low	

Add your score across all 20 items, then check it against the tier table above. That number is your starting point — not a verdict.

4. Measurement & Attribution

Section 1 named the measurement gap. This section explains what's actually causing it, and how to close it.

Standard GA4 configurations weren't built to distinguish an AI-referred session from a generic referral or direct visit. When someone arrives at your site after an AI system named you in an

answer, GA4 typically has no clean way to label that visit as AI-referred unless you've deliberately configured it to. The traffic still gets counted — it just gets counted as something else, usually direct or unassigned referral, which means the revenue and conversions attached to it disappear into a bucket that doesn't reflect where they actually came from.

Conversion behavior compounds the problem. Where AI-referred traffic is visible, it converts at rates that diverge from organic search — sometimes higher, since a visitor arriving via an AI recommendation often arrives pre-qualified by the synthesis that got them there; sometimes lower, depending on where in their research process that recommendation happened. Blending it into a single, undifferentiated organic conversion rate hides which is actually true for your brand.

The practical result: brands making decisions from blended dashboards are frequently optimizing against numbers that don't represent the channel they're trying to grow.

Callout: Minimum Viable AI Attribution Setup

- Segment AI referrers — chatgpt.com, perplexity.ai, claude.ai, gemini.google.com, and similar — as their own channel, not folded into "referral" or "direct"
- Tag AI-landing pages so you can see exactly which content AI systems are sending visitors to, not just that traffic arrived
- Monitor citation appearances on a recurring cadence — monthly at minimum — using the same branded and category prompts from your Part A audit
- Report AI-referred conversion rate and revenue as their own line items, never blended into overall organic performance

Once that foundation is in place, the reporting itself should live in one view, not scattered across GA4, a schema checker, and a rank tracker separately. FlyRank's own free AI Visibility Insights product is structured around exactly this — one report covering readiness, traffic, revenue, conversion, and next actions together:

Callout: AI Visibility Reporting Template

Report Section	What It Tracks	Why It's Tracked Separately
LLM Readiness	Current state of schema, llms.txt, robots.txt, and crawl access	This is your Pillar 2 audit score — a leading indicator, not a lagging one
AI-Platform Traffic	Sessions referred from	Separates AI-referred visitors

Report Section	What It Tracks	Why It's Tracked Separately
	ChatGPT, Perplexity, Claude, Gemini, and similar	from organic, direct, and paid
AI-Referred Revenue	Revenue attributed specifically to AI-referred sessions	Surfaces the under-attribution problem directly, instead of leaving it hidden
AI Conversion Rate	Conversion rate for AI-referred sessions, benchmarked against organic	Confirms whether AI traffic converts differently, and by how much
AI Landing Pages	The specific pages AI systems are actually sending traffic to	Shows which content is being cited and clicked, not just published
Growth Sources	Which platform, query type, or content piece is driving the most movement	Prioritizes where to invest next, instead of guessing
Action Items	Specific fixes tied back to gaps identified in your Section 3 audit	Keeps reporting connected to remediation, not just observation

FlyRank's own client results reflect exactly what this structure is built to catch. Becoming the cited answer across both Google and AI systems like ChatGPT tends to show up first in citation tracking and AI-referred traffic — well before it moves a blended revenue number enough to notice. Brands watching only the blended number typically find out months later than brands watching the AI-specific lines.

5. From Framework to Standing Advantage

The Bridge. You now understand how AI systems evaluate a brand — through mentions,

consistent facts, and third-party corroboration. What the checklist can't tell you is that none of these signals hold still. Mentions decay, schema drifts out of sync with product changes, and competitors earn new citations every week — which makes this a continuous monitoring and correction problem, not a one-time fix.

The Offer. Running that continuously, across every channel, for every fact your brand puts into the world, is exactly what FlyRank does for clients every day.

The Action.

- **Primary:** Visit flyrank.ai and submit the form to start your free AI Visibility audit.
- **Fallback:** Prefer to talk it through first? Email the team directly at growth@flyrank.com.